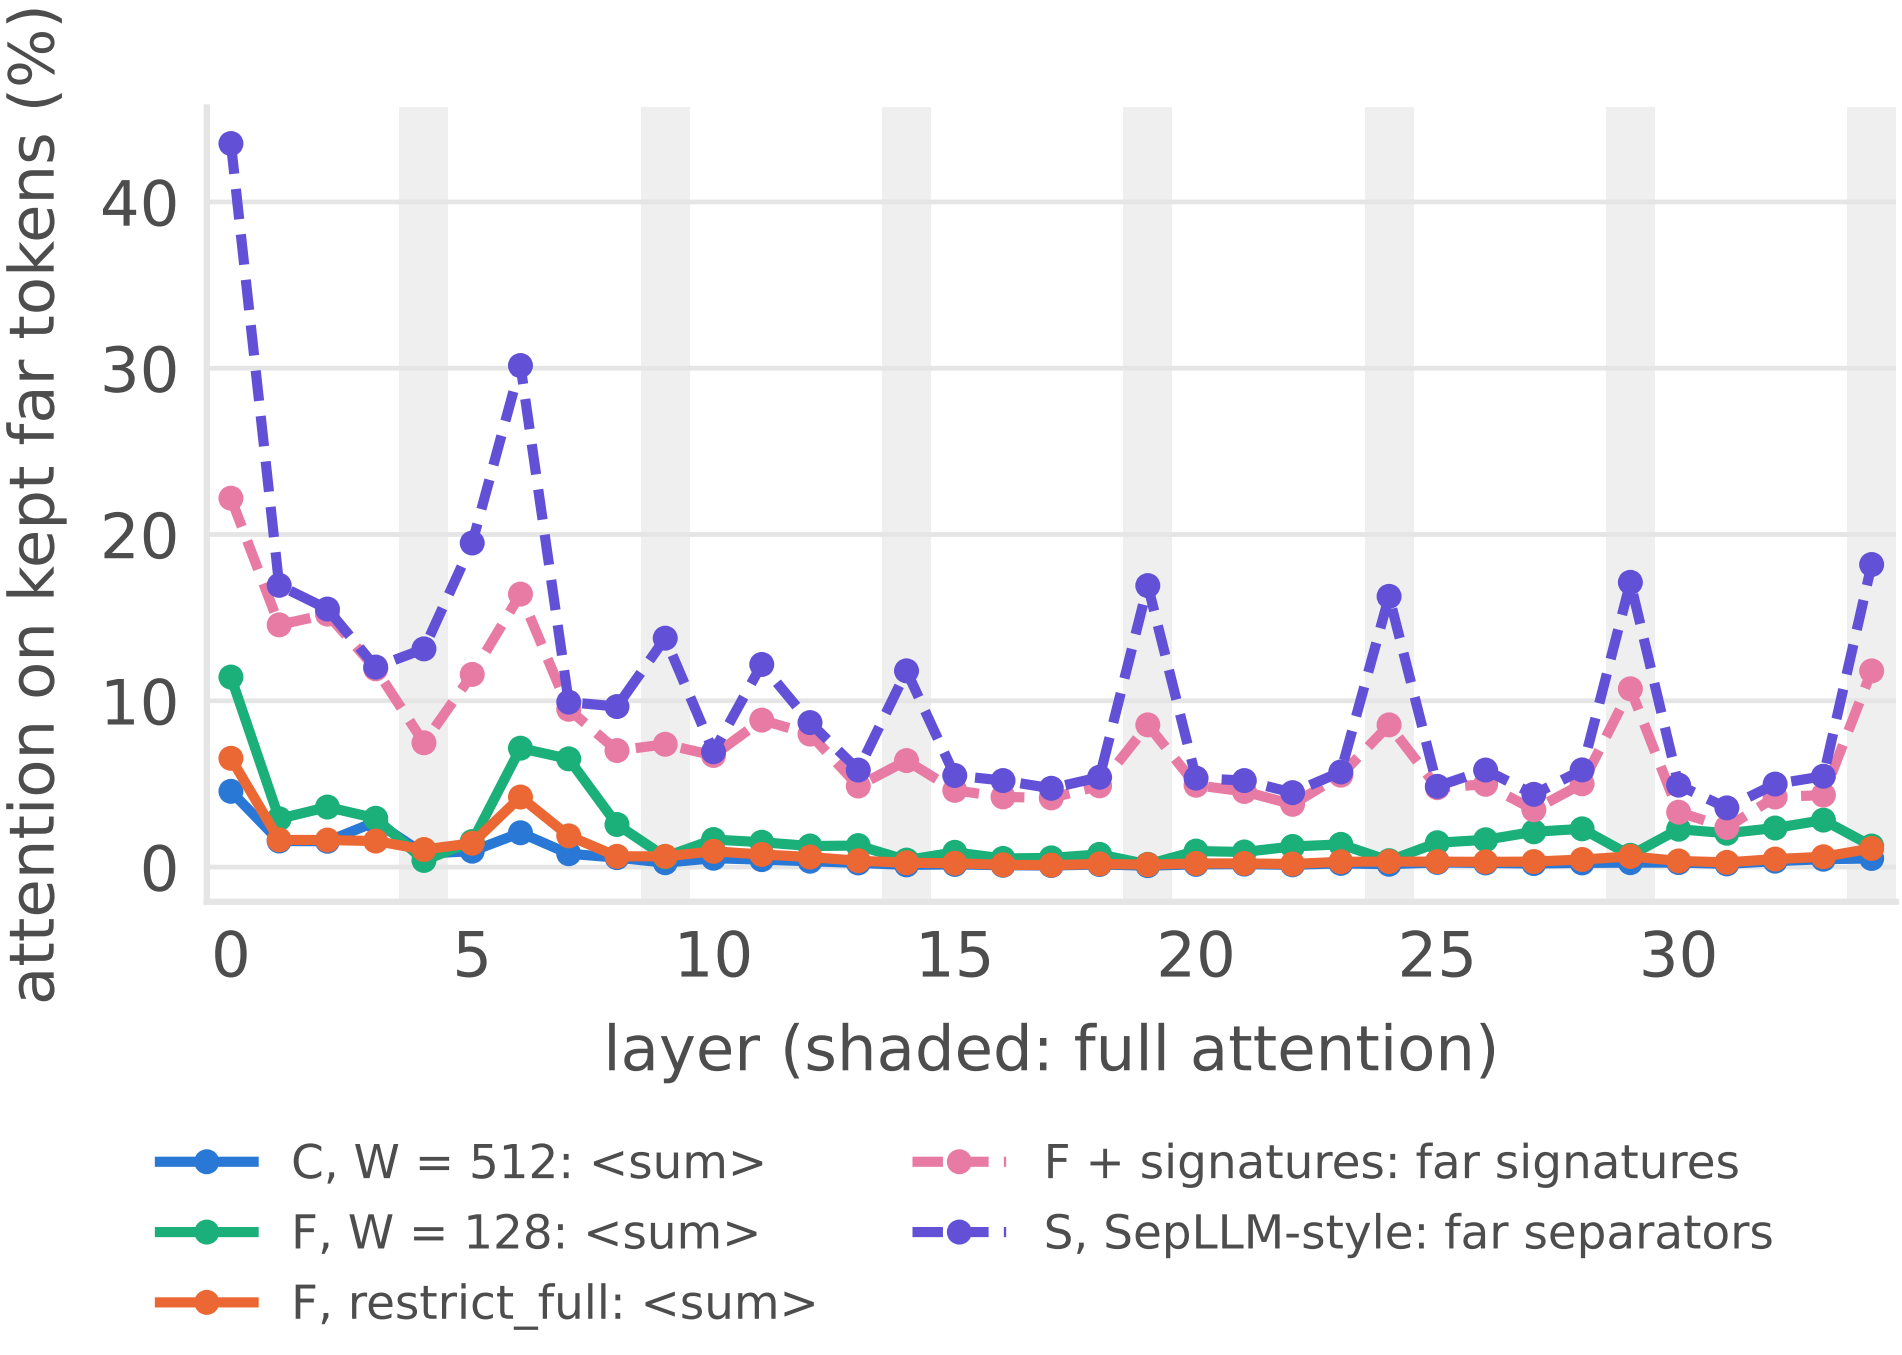
attention on kept far tokens (%)
0
10
20
30
40
0
5
10
15
20
25
30
layer (shaded: full attention)
C, W = 512: <sum>
F, W = 128: <sum>
F, restrict_full: <sum>
F + signatures: far signatures
S, SepLLM-style: far separators